

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense, LSTM
import tensorflow
from sklearn.preprocessing import MinMaxScaler

%matplotlib inline

In [125]: df = pd.read_csv('TSLA.csv', index_col=0)
df.index = pd.to_datetime(df.index)

#create classes, and shift data back one (yesterdays parameters should predict today's outcome)
df['Class'] = df['Close'] - df['Open'] > 0).mul(1).shift(-1)
df.dropna(inplace=True)

#create lag in dataframe which will tell the nth last class
for i in range(1, 6):
    df['lag{}'.format(i)] = df['Class'].shift(i)

df['100mav'] = df['Close'].shift(1).rolling(100).mean()
df['50mav'] = df['Close'].shift(1).rolling(50).mean()
df['30mav'] = df['Close'].shift(1).rolling(30).mean()
df['10mav'] = df['Close'].shift(1).rolling(10).mean()
df['30d_high'] = df['Close'].shift(1).rolling(30).max()
df['20d_high'] = df['Close'].shift(1).rolling(20).max()
df['10d_high'] = df['Close'].shift(1).rolling(10).max()
df['5d_high'] = df['Close'].shift(1).rolling(5).max()
# df.head(20)

df.dropna(inplace=True)

In [126]: # Dataframe with only close column
data_close = df.filter(['Close'])
# Convert the dataframe to a numpy array
dataset_close = data_close.values
# Get the number of rows to train the model on
training_data_len = int(np.ceil(len(dataset_close) * .66))

In [127]: # Scale the data
close_scaler = MinMaxScaler(feature_range=(0, 1))
close_scaled_data = close_scaler.fit_transform(dataset_close)
close_scaled_data

Out[127]: array([[0.00899965],
[0.01022789],
[0.01291886],
...,
[0.54995031],
[0.53629453],
[0.56071418]])

In [128]: # Create training data set
# Create scaled training data set
train_data = close_scaled_data[0:int(training_data_len), :]
# Split data into x-train and y-train data sets
x_train = []
y_train = []

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])
    if i <= 61:
        print(x_train)
        print(y_train)
        print()

# Convert the x_train and y_train to numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Reshape the data
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

[array([[0.00899965, 0.01022789, 0.01291886, 0.01422526, 0.01523019,
0.0150627 , 0.01395728, 0.01507386, 0.01397961, 0.01174645,
0.01078619, 0.00946862, 0.01086435, 0.01176878, 0.01141147,
0.01081968, 0.0097366 , 0.0074811 , 0.00867585, 0.01002691,
0.01064103, 0.01182067, 0.01164595, 0.01205909, 0.00922297,
0.00415369, 0.00511395, 0.00658784, 0.00521444, 0.00535959,
0.00534843, 0.00540426, 0.00558291, 0.00675532, 0.00715729,
0.00739178, 0.00572807, 0.00572807, 0.0049018 , 0.004377 , 0.0025418,
0.00425418, 0.00245648, 0.0008821 , 0.00135107, 0.00297011,
0.00318226, 0.00236042, 0.00345024, 0.00243415, 0.00253464,
0.00232249, 0.00235599, 0.00200985, 0.00182003, 0.00138456,
0.00297011, 0.00154088, 0.00155205, 0.00158555, 0.00139573]),
[0.0011277484653344112]

[array([[0.00899965, 0.01022789, 0.01291886, 0.01422526, 0.01523019,
0.0150627 , 0.01395728, 0.01507386, 0.01397961, 0.01174645,
0.01078619, 0.00946862, 0.01086435, 0.01176878, 0.01141147,
0.01081968, 0.0097366 , 0.0074811 , 0.00867585, 0.01002691,
0.01064103, 0.01182067, 0.01164595, 0.01205909, 0.00922297, 0.00415369,
0.00511395, 0.00658784, 0.00521444, 0.00535959, 0.00534843,
0.00540426, 0.00558291, 0.00675532, 0.00715729, 0.00739178,
0.00572807, 0.00572807, 0.0049018 , 0.004377 , 0.0025418,
0.00425418, 0.00245648, 0.0008821 , 0.00135107, 0.00297011, 0.00318226,
0.00236042, 0.00345024, 0.00243415, 0.00253464, 0.00232249,
0.00235599, 0.00200985, 0.00182003, 0.00138456, 0.00297011,
0.00154088, 0.00155205, 0.00158555, 0.00139573]), array([[0.01022789, 0.01291886, 0.01422526, 0.01523019,
0.0150627 ,
0.01395728, 0.01507386, 0.01397961, 0.01174645, 0.01078619,
0.00946862, 0.01086435, 0.01176878, 0.01141147, 0.01081968,
0.0097366 , 0.0074811 , 0.00867585, 0.01002691, 0.01064103,
0.01182067, 0.01164595, 0.01205909, 0.00922297, 0.00415369,
0.00511395, 0.00658784, 0.00521444, 0.00535959, 0.00534843,
0.00540426, 0.00558291, 0.00675532, 0.00715729, 0.00739178,
0.00572807, 0.00572807, 0.0049018 , 0.004377 , 0.0025418,
0.00425418, 0.00245648, 0.0008821 , 0.00135107, 0.00297011,
0.00318226, 0.00236042, 0.00345024, 0.00243415, 0.00253464,
0.00232249, 0.00235599, 0.00200985, 0.00182003, 0.00138456,
0.00297011, 0.00154088, 0.00155205, 0.00158555, 0.00139573]),
[0.0011277484653344112, 0.003238088493196279]

In [129]: # Build LSTM model
model = Sequential()
model.add(LSTM(128, return_sequences=True, input_shape=(x_train.shape[1], 1)))
model.add(LSTM(64, return_sequences=False))
model.add(Dense(25))
model.add(Dense(1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Train the model
model.fit(x_train, y_train, batch_size=1, epochs=1)

1495/1495 [=====] - 33s 20ms/step - loss: 3.1415e-04

Out[129]: <keras.callbacks.History at 0x20484777890>

In [130]: # Create the testing data set
# Create a new array containing scaled values from index 1543 to 2002
test_data = close_scaled_data[training_data_len - 60: , :]
# Create the data sets x_test and y_test
x_test = []
y_test = dataset_close[training_data_len:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

# Convert the data to a numpy array
x_test, y_test = np.array(x_test)

# Reshape the data
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

# Get the models predicted price values
predictions = model.predict(x_test)
predictions = close_scaler.inverse_transform(predictions)

# Get the root mean squared error (RMSE)
rmse = np.sqrt(np.mean(((predictions - y_test) ** 2)))
rmse

26/26 [=====] - 2s 21ms/step

Out[130]: 37.11085253419023

In [131]: # Plot the data
train = data_close[:training_data_len]
close_valid = data_close[training_data_len:]
close_valid['Close Predictions'] = predictions
# Visualize the data
plt.figure(figsize=(16,6))
plt.title('Model')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train['Close'])
plt.plot(close_valid[['Close', 'Close Predictions']])
plt.legend(['Train', 'Val', 'Close Predictions'], loc='lower right')
plt.show()

C:\Users\Brandon\AppData\Local\Temp\ipykernel_2364\2370582738.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
close_valid['Close Predictions'] = predictions

Model

Close Price USD ($)

800
600
400
200
0
2011 2012 2013 2014 2015 2016 2017 2018 2019 2020
Train
Val
Close Predictions

In [132]: close_valid

Out[132]: Close Close Predictions

Date

2017-01-25 254.470001 236.691345
2017-01-26 252.509995 239.623901
2017-01-27 252.949997 241.257019
2017-01-30 250.630005 242.027954
2017-01-31 251.929993 241.742233
... ... ...
2020-03-25 539.250000 401.491364
2020-03-26 528.159973 433.839844
2020-03-27 514.359985 460.170105
2020-03-30 502.130005 474.709534
2020-03-31 524.000000 477.820374

801 rows x 2 columns

Now that we made a model for the closing data, let's make the same model but for open.
```

```
In [133]: # Dataframe with only Open column
data_open = df.filter(['Open'])
# Convert the dataframe to a numpy array
dataset_open = data_open.values
# Get the number of rows to train the model on
training_data_len = int(np.ceil(len(dataset_open) * .66))

In [134]: # Scale the data
open_scaler = MinMaxScaler(feature_range=(0, 1))
open_scaled_data = open_scaler.fit_transform(dataset_open)
open_scaled_data

Out[134]: array([[0.00985894],
[0.00929335],
[0.01085703],
...,
[0.53588697],
[0.54172028],
[0.53172825]])

In [135]: # Create training data set
# Create scaled training data set
train_data = open_scaled_data[0:int(training_data_len), :]
# Split data into x-train and y-train data sets
x_train = []
y_train = []

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])
    if i <= 61:
        print(x_train)
        print(y_train)
        print()

# Convert the x_train and y_train to numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Reshape the data
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

[array([[0.00985894, 0.00929335, 0.01085703, 0.01276449, 0.01496603 ,
0.01532626, 0.01511556, 0.01326354, 0.01562569, 0.01356297,
0.01134498, 0.01061305, 0.00965932, 0.01186621, 0.01189948, 0.01138934,
0.01138934, 0.01093466, 0.00943752, 0.00764095, 0.00911591,
0.01060196, 0.01093466, 0.0111121 , 0.01161114, 0.01051324,
0.00692011, 0.0045136 , 0.00582221, 0.00565523, 0.00597963,
0.0056115 , 0.00541188, 0.00521226, 0.00560041, 0.00689793,
0.00708646, 0.00755223, 0.00580003, 0.00574458, 0.00484629,
0.00410327, 0.00387038, 0.00249523, 0.00148605, 0.00194074,
0.00318281, 0.0024934 , 0.00328261, 0.00343787, 0.00251741,
0.00208575, 0.0026394 , 0.00226234, 0.00184093, 0.00164131,
0.00221798, 0.00260613, 0.00164131, 0.00163022, 0.00206272]),
[0.0013640592902153213]

[array([[0.00985894, 0.00929335, 0.01085703, 0.01276449, 0.01496603 ,
0.01532626, 0.01511556, 0.01326354, 0.01562569, 0.01356297,
0.01134498, 0.01061305, 0.00965932, 0.01186621, 0.01189948, 0.01138934,
0.01138934, 0.01093466, 0.00943752, 0.00764095, 0.00911591,
0.01060196, 0.01093466, 0.0111121 , 0.01161114, 0.01051324,
0.00692011, 0.0045136 , 0.00582221, 0.00565523, 0.00597963,
0.0056115 , 0.00541188, 0.00521226, 0.00560041, 0.00689793,
0.00708646, 0.00755223, 0.00580003, 0.00574458, 0.00484629,
0.00410327, 0.00387038, 0.00249523, 0.00148605, 0.00194074,
0.00318281, 0.0024934 , 0.00328261, 0.00343787, 0.00251741,
0.00208575, 0.0026394 , 0.00226234, 0.00184093, 0.00164131,
0.00221798, 0.00260613, 0.00164131, 0.00163022, 0.00206272]),
array([[0.00929335, 0.01085703, 0.01276449, 0.01496603
, 0.01532626,
0.01511556, 0.01326354, 0.01562569, 0.01356297, 0.01134498,
0.01061305, 0.00965932, 0.01186621, 0.01189948, 0.01138934,
0.01093466, 0.00943752, 0.00764095, 0.00911591, 0.01060196,
0.01093466, 0.0111121 , 0.01161114, 0.01051324, 0.00692011,
0.0045136 , 0.00582221, 0.00565523, 0.00597963, 0.0056115 ,
0.00541188, 0.00521226, 0.00560041, 0.00689793, 0.00708646,
0.00755223, 0.00580003, 0.00574458, 0.00484629, 0.00410327,
0.00387038, 0.00249523, 0.00148605, 0.00194074, 0.00318281,
0.0024934 , 0.00328261, 0.00343787, 0.00251741, 0.00208575,
0.0026394 , 0.00226234, 0.00184093, 0.00164131, 0.00221798,
0.00260613, 0.00164131, 0.00163022, 0.00206272, 0.00136406]),
[0.0013640592902153213, 0.0014638687129368104]

In [136]: # Build LSTM model
model = Sequential()
model.add(LSTM(128, return_sequences=True, input_shape=(x_train.shape[1], 1)))
model.add(LSTM(64, return_sequences=False))
model.add(Dense(25))
model.add(Dense(1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Train the model
model.fit(x_train, y_train, batch_size=1, epochs=1)

1495/1495 [=====] - 33s 20ms/step - loss: 2.8712e-04

Out[136]: <keras.callbacks.History at 0x2048a1f53d0>

In [137]: # Create the testing data set
# Create a new array containing scaled values from index 1543 to 2002
test_data = open_scaled_data[training_data_len - 60: , :]
# Create the data sets x_test and y_test
x_test = []
y_test = dataset_open[training_data_len:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

# Convert the data to a numpy array
x_test, y_test = np.array(x_test)

# Reshape the data
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

# Get the models predicted price values
predictions = model.predict(x_test)
predictions = open_scaler.inverse_transform(predictions)

# Get the root mean squared error (RMSE)
rmse = np.sqrt(np.mean(((predictions - y_test) ** 2)))
rmse

26/26 [=====] - 2s 31ms/step

Out[137]: 29.316242915296616

In [138]: # Plot the data
train = data_open[:training_data_len]
open_valid = data_open[training_data_len:]
open_valid['Open Predictions'] = predictions
# Visualize the data
plt.figure(figsize=(16,6))
plt.title('Model')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Open Price USD ($)', fontsize=18)
plt.plot(train['Open'])
plt.plot(open_valid[['Open', 'Open Predictions']])
plt.legend(['Train', 'Val', 'Open Predictions'], loc='lower right')
plt.show()

C:\Users\Brandon\AppData\Local\Temp\ipykernel_2364\2352746716.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
open_valid['Open Predictions'] = predictions

Model

Open Price USD ($)

800
600
400
200
0
2011 2012 2013 2014 2015 2016 2017 2018 2019 2020
Train
Val
Open Predictions

In [140]: open_valid

Out[140]: Open Open Predictions

Date

2017-01-25 257.309998 258.268036
2017-01-26 254.289993 261.459412
2017-01-27 251.380005 264.079285
2017-01-30 252.529999 265.595978
2017-01-31 249.240005 266.456696
... ... ...
2020-03-25 545.250000 417.985779
2020-03-26 547.390015 437.734467
2020-03-27 505.000000 462.927582
2020-03-30 510.260010 480.852173
2020-03-31 501.250000 493.060364

801 rows x 2 columns

In [141]: close_valid

Out[141]: Close Close Predictions

Date

2017-01-25 254.470001 236.691345
2017-01-26 252.509995 239.623901
2017-01-27 252.949997 241.257019
2017-01-30 250.630005 242.027954
2017-01-31 251.929993 241.742233
... ... ...
2020-03-24 378.555542 414.576019
2020-03-25 401.491364 433.839844
2020-03-26 547.390015 437.734467
2020-03-27 505.000000 462.927582
2020-03-30 510.260010 480.852173
2020-03-31 501.250000 493.060364

800 rows x 3 columns

Does not give desired results, since our model's open and close predictions never overlap. Instead, we can assume that the previous day close is the next day's open.
```

```
In [164]: close_valid['Class Prediction'] = close_valid['Close Predictions'].gt(close_valid['Close Predictions'].shift())
close_valid.head(1)

C:\Users\Brandon\AppData\Local\Temp\ipykernel_2364\2352746716.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
close_valid['Class Prediction'] = close_valid['Close Predictions'].gt(close_valid['Close Predictions'].shift())

Out[164]: Close Predictions Open Predictions Class

Date

2017-01-25 236.691345 258.268036 0.0
2017-01-26 239.623901 261.459412 0.0
2017-01-27 241.257019 264.079285 0.0
2017-01-30 242.027954 265.595978 0.0
2017-01-31 241.742233 266.456696 0.0
... ... ...
2020-03-24 378.555542 414.576019 0.0
2020-03-25 401.491364 433.839844 0.0
2020-03-26 547.390015 437.734467 0.0
2020-03-27 505.000000 462.927582 0.0
2020-03-30 510.260010 480.852173 0.0
2020-03-31 501.250000 493.060364 0.0

800 rows x 3 columns

These results are more desirable, let's now compare it to the actual values
```

```
In [182]: y_test = df.Class.iloc[training_data_len:]
y_pred = close_valid['Class Prediction']
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
print('True negative:', tn)
print('False Positive:', fp)
print('False Negative:', fn)
print('True Positive:', tp)

print('Accuracy Rate:', (tp + tn)/(fn+fp+tn+tp))

True negative: 176
False Positive: 213
False Negative: 185
True Positive: 227
Accuracy Rate: 0.5031210986267166

Unfortunately, our model is only a little better than a simple coin toss.
```

